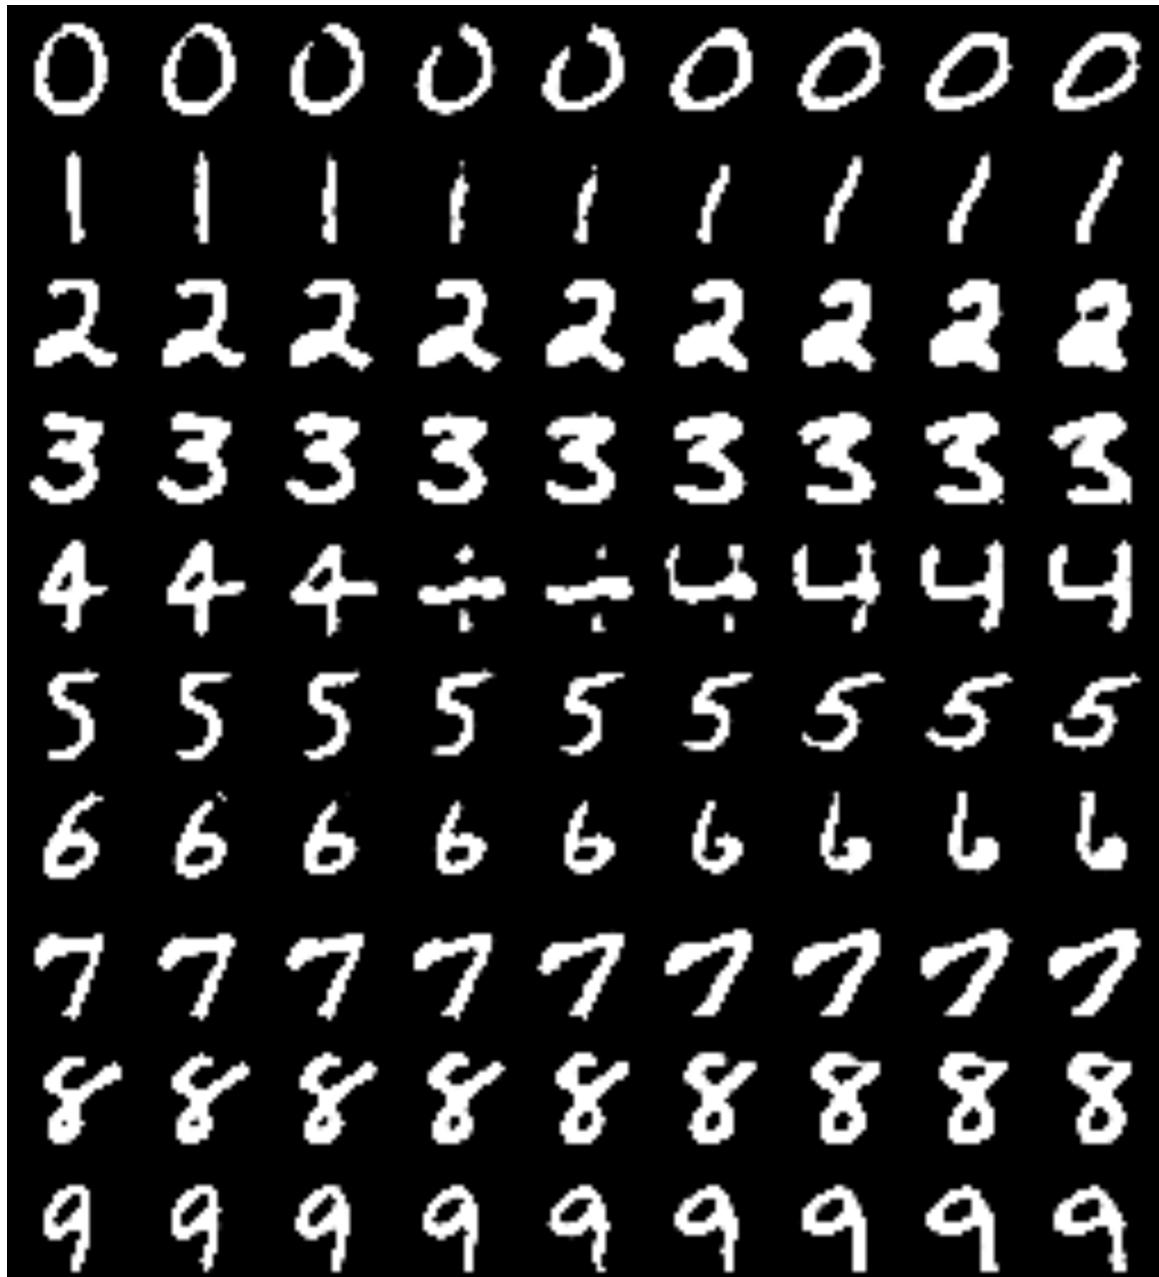
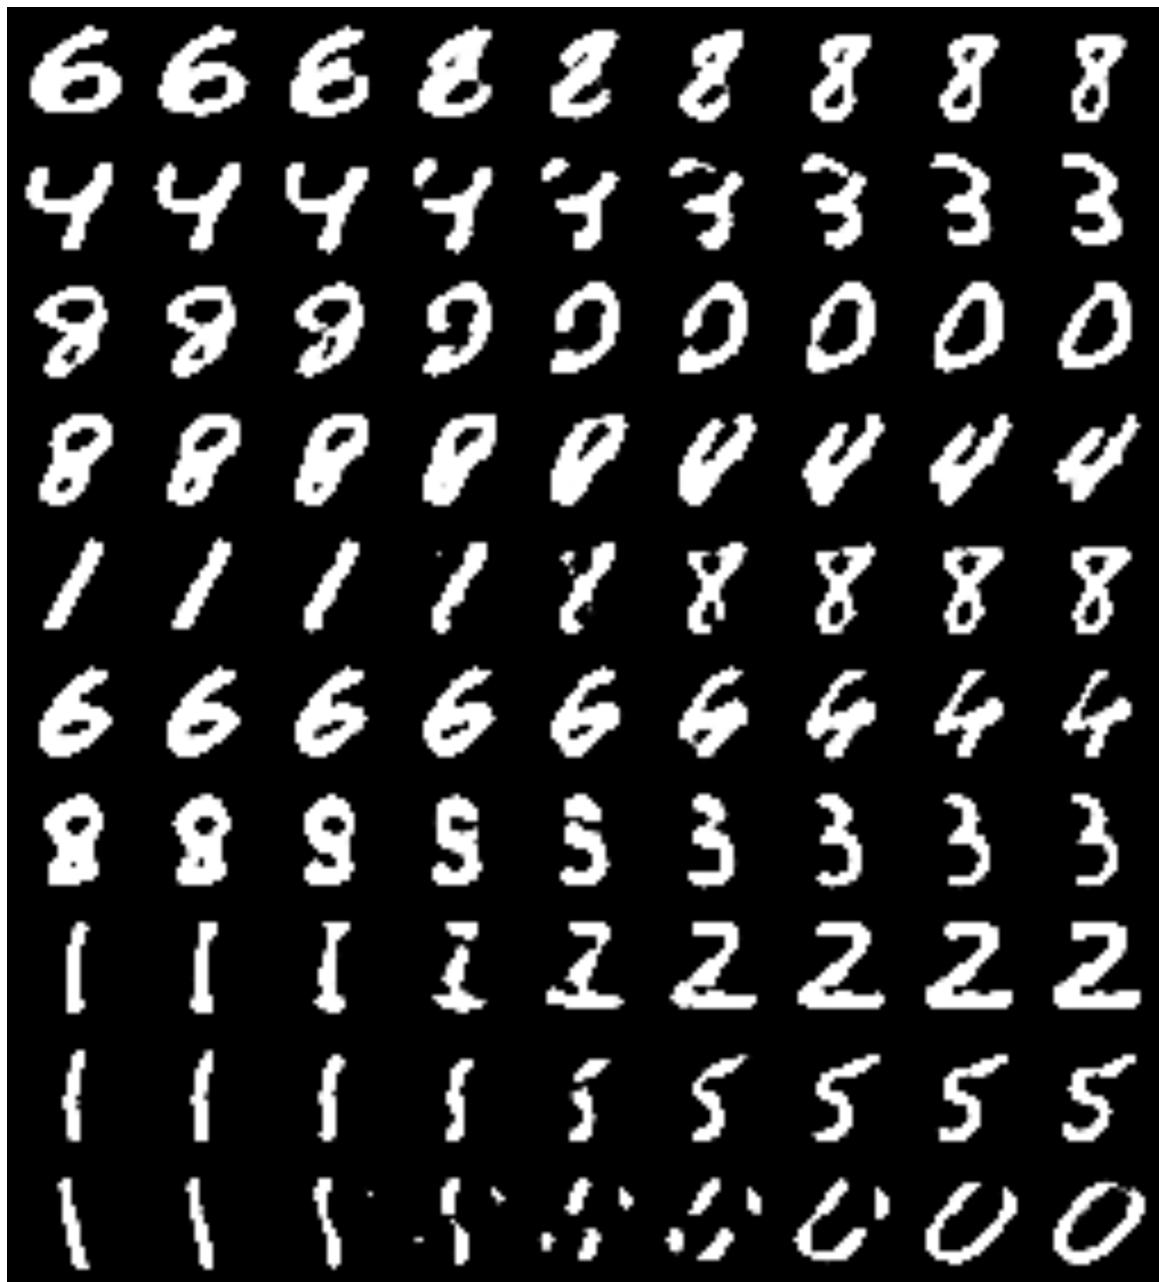


Homework 9: Variational Autoencoders

1. (45 points) Page 1: Same digit interpolates (Make your figure large!)



2. (45 points) Page 2: Different digit interpolates (Make your figure large!)



3. (10 points) Page 3+: Your code.

```
[2] !mkdir hw9_data

[4] from torchvision import datasets, transforms
    ## YOUR CODE HERE ##
    transformations = transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.1307,), (0.3081,)))
    ])
mnist_train = datasets.MNIST(root = 'hw9_data', transform = transformations, train = True, download = True)
mnist_test = datasets.MNIST(root = 'hw9_data', transform = transformations, train = False, download = True)

[5] from torch.utils.data import DataLoader
    ## YOUR CODE HERE ##
train_loader = DataLoader(dataset = mnist_train, batch_size = 32, shuffle = True)
test_loader = DataLoader(dataset = mnist_test, batch_size = 32, shuffle = True)

[14] import os
    import torch
    import torch.utils.data
    from torch import nn, optim
    from torch.autograd import Variable
    from torch.nn import functional as F
    from torchvision import datasets, transforms
    from torchvision.utils import save_image

    # changed configuration to this instead of argparse for easier interaction
    CUDA = False
    SEED = 1
    BATCH_SIZE = 128
    LOG_INTERVAL = 10
    EPOCHS = 10

[14] class VAE(nn.Module):
    def __init__(self):
        super(VAE, self).__init__()

        # ENCODER
        # 28 x 28 pixels = 784 input pixels, 400 outputs
        self.fc1 = nn.Linear(784, 400)
        # rectified linear unit layer from 400 to 400
        # max(0, x)
        self.relu = nn.ReLU()
        self.fc21 = nn.Linear(400, ZDIMS) # mu layer
        self.fc22 = nn.Linear(400, ZDIMS) # logvariance layer
        # this last layer bottlenecks through ZDIMS connections

        # DECODER
        # from bottleneck to hidden 400
        self.fc3 = nn.Linear(ZDIMS, 400)
        # from hidden 400 to 784 outputs
        self.fc4 = nn.Linear(400, 784)
        self.sigmoid = nn.Sigmoid()

    def encode(self, x: Variable) -> (Variable, Variable):
        """Input vector x -> fully connected 1 -> ReLU -> (fully connected
        21, fully connected 22)

        Parameters
        -----
        x : [128, 784] matrix; 128 digits of 28x28 pixels each

        Returns
        -----
        (mu, logvar) : ZDIMS mean units one for each latent dimension, ZDIMS
        variance units one for each latent dimension

        """

```

```

[14]     # h1 is [128, 400]
    h1 = self.relu(self.fc1(x)) # type: Variable
    return self.fc21(h1), self.fc22(h1)

def reparameterize(self, mu: Variable, logvar: Variable) -> Variable:
    """THE REPARAMETERIZATION IDEA:

    For each training sample (we get 128 batched at a time)

    - take the current learned mu, stddev for each of the ZDIMS
      dimensions and draw a random sample from that distribution
    - the whole network is trained so that these randomly drawn
      samples decode to output that looks like the input
    - which will mean that the std, mu will be learned
      *distributions* that correctly encode the inputs
    - due to the additional KLD term (see loss_function() below)
      the distribution will tend to unit Gaussians

    Parameters
    -----
    mu : [128, ZDIMS] mean matrix
    logvar : [128, ZDIMS] variance matrix

    Returns
    -----
    During training random sample from the learned ZDIMS-dimensional
    normal distribution; during inference its mean.

    """
    if self.training:
        # multiply log variance with 0.5, then in-place exponent
        # yielding the standard deviation
        std = logvar.mul(0.5).exp_() # type: Variable
        # - std.data is the [128,ZDIMS] tensor that is wrapped by std
        # - so eps is [128,ZDIMS] with all elements drawn from a mean 0
        #   and stddev 1 normal distribution that is 128 samples
        #   of random ZDIMS-float vectors
        eps = Variable(std.data.new(std.size()).normal_())
        # - sample from a normal distribution with standard
        #   deviation = std and mean = mu by multiplying mean 0
        #   stddev 1 sample with desired std and mu, see
        #   https://stats.stackexchange.com/a/16338
        # - so we have 128 sets (the batch) of random ZDIMS-float
        #   vectors sampled from normal distribution with learned
        #   std and mu for the current input
        return eps.mul(std).add_(mu)

    else:
        # During inference, we simply spit out the mean of the
        # learned distribution for the current input. We could
        # use a random sample from the distribution, but mu of
        # course has the highest probability.
        return mu

def decode(self, z: Variable) -> Variable:
    h3 = self.relu(self.fc3(z))
    return self.sigmoid(self.fc4(h3))

def forward(self, x: Variable) -> (Variable, Variable, Variable):
    mu, logvar = self.encode(x.view(-1, 784))
    z = self.reparameterize(mu, logvar)
    return self.decode(z), z, mu, logvar

model = VAE()
if CUDA:
    model.cuda()

```

```
[14] def loss_function(recon_x, x, mu, logvar) -> Variable:
    # how well do input x and output recon_x agree?
    BCE = F.binary_cross_entropy(recon_x, x.view(-1, 784))

    # KLD is Kullback-Leibler divergence -- how much does one learned
    # distribution deviate from another, in this specific case the
    # learned distribution from the unit Gaussian

    # see Appendix B from VAE paper:
    # Kingma and Welling. Auto-Encoding Variational Bayes. ICLR, 2014
    # https://arxiv.org/abs/1312.6114
    # - D_{KL} = 0.5 * sum(1 + log(sigma^2) - mu^2 - sigma^2)
    # note the negative D_{KL} in appendix B of the paper
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    # Normalise by same number of elements as in reconstruction
    KLD /= BATCH_SIZE * 784

    # BCE tries to make our reconstruction as accurate as possible
    # KLD tries to push the distributions as close as possible to unit Gaussian
    return BCE + KLD

# Dr Diederik Kingma: as if VAEs weren't enough, he also gave us Adam!
optimizer = optim.Adam(model.parameters(), lr=1e-3)

def train(epoch):
    # toggle model to train mode
    model.train()
    train_loss = 0
    # in the case of MNIST, len(train_loader.dataset) is 60000
    # each `data` is of BATCH_SIZE samples and has shape [128, 1, 28, 28]
    for batch_idx, (data, _) in enumerate(train_loader):
        data = Variable(data)
        if CUDA:
            data = data.cuda()
        optimizer.zero_grad()

[14]     # push whole batch of data through VAE.forward() to get recon_loss
        recon_batch, code, mu, logvar = model(data)
        # calculate scalar loss
        loss = loss_function(recon_batch, data, mu, logvar)
        # calculate the gradient of the loss w.r.t. the graph leaves
        # i.e. input variables -- by the power of pytorch!
        loss.backward()
        train_loss += loss.item()
        optimizer.step()
        if batch_idx % LOG_INTERVAL == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader),
                loss.item() / len(data)))
    print('====> Epoch: {} Average loss: {:.4f}'.format(
        epoch, train_loss / len(train_loader.dataset)))

def test(epoch):
    # toggle model to test / inference mode
    model.eval()
    test_loss = 0

    # each data is of BATCH_SIZE (default 128) samples
    for i, (data, _) in enumerate(test_loader):
        if CUDA:
            # make sure this lives on the GPU
            data = data.cuda()

            # we're only going to infer, so no autograd at all required: volatile=True
            data = Variable(data, volatile=True)
        recon_batch, mu, logvar = model(data)
        test_loss += loss_function(recon_batch, data, mu, logvar).item()
        if i == 0:
            n = min(data.size(0), 8)
            # for the first 128 batch of the epoch, show the first 8 input digits
            # with right below them the reconstructed output digits
```

```
[14]         comparison = torch.cat([data[:n],
                           recon_batch.view(BATCH_SIZE, 1, 28, 28)[:n]])
        save_image(comparison.data.cpu(),
                   'results/reconstruction_' + str(epoch) + '.png', nrow=n)

    test_loss /= len(test_loader.dataset)
    print('====> Test set loss: {:.4f}'.format(test_loss))

    for epoch in range(1, EPOCHS + 1):
        train(epoch)
        #     test(epoch)

        # 64 sets of random ZDIMS-float vectors, i.e. 64 locations / MNIST
        # digits in latent space
        sample = Variable(torch.randn(64, ZDIMS))
        if CUDA:
            sample = sample.cuda()
        sample = model.decode(sample).cpu()
```

▼ Part 4: Visualizing the VAE output (90 points)

We will look at how well the codes produced by the VAE can be interpolated. **For this section we will only use the MNIST test set.**

To create an interpolation between two images A and B, we encode both these images and generate a sample code for each of them. We now consider 7 equally spaced points in between these two sample codes giving us a total of 9 points including the samples. We then decode these images to get interpolated images in between A and B.

Complete the interpolation function below that takes a pair of images A and B and returns 9 images. (You are free to use any data structure you want to return these images)

```
[18] import matplotlib.pyplot as plt
from torchvision import utils
%matplotlib inline
import numpy as np

def create_interpolates(A, B, vae):
    ## YOUR CODE HERE ##
    mu_a, logvar_a = model.encode(A)
    mu_b, logvar_b = model.encode(B)
    code_a = model.reparameterize(mu_a, logvar_a)
    code_b = model.reparameterize(mu_b, logvar_b)
    interval = (code_b - code_a)/8
    index = 1
    for i in range(8):
        code = code_a + interval * i
        image = model.decode(code).view(28, 28)
        if index:
            interpolates = image
            index = 0
        else:
            interpolates = torch.cat((interpolates, image))
    interpolates = torch.cat((interpolates, model.decode(code_b).view(28, 28)))
    interpolates = interpolates.view(9, 1, 28, 28)
    grid = utils.make_grid(interpolates, nrow = 1)
    show(grid)
```

```

def show(img):
    npimg = img.detach().numpy()
    plt.imshow(np.transpose(npimg, (1,2,0)), interpolation='nearest')

similar_pairs = {}
for _, (x, y) in enumerate(test_loader):
    for i in range(len(y)):
        if y[i].item() not in similar_pairs:
            similar_pairs[y[i].item()] = []
        if len(similar_pairs[y[i].item()]) < 2:
            similar_pairs[y[i].item()].append(x[i])

done = True
for i in range(10):
    if i not in similar_pairs or len(similar_pairs[i]) < 2:
        done = False

if done:
    break

# similar_pairs[i] contains two images indexed at 0 and 1 that have images of the digit i

## YOUR CODE HERE ##
interpolates = []
index = 1
for i in range(10):
    digits = similar_pairs[i]
    digit1 = digits[0].view(-1, 28*28)
    digit2 = digits[1].view(-1, 28*28)
    interpolate = create_interpolates(digit1, digit2, model)
    if index:
        interpolates = interpolate
        index = 0
    else:
        interpolates = torch.cat((interpolates, interpolate))

```

```

transform = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize(size=800),
    transforms.ToTensor()
])
interpolates = [transform(interpolates_) for interpolates_ in interpolates]
grid = utils.make_grid(interpolates, nrow = 9)
utils.save_image(grid, 'same_digit.png')
show(grid)

```

```

[34] random_pairs = {}
for _, (x, y) in enumerate(test_loader):
    # Make sure the batch size is greater than 20
    for i in range(10):
        random_pairs[i] = []
        random_pairs[i].append(x[2*i])
        random_pairs[i].append(x[2*i+1])
    break

# random_pairs[i] contains two images indexed at 0 and 1 that are chosen at random

## YOUR CODE HERE ##
interpolates = []
index = 1
for i in range(10):
    digits = random_pairs[i]
    digit1 = digits[0].view(-1, 28*28)
    digit2 = digits[1].view(-1, 28*28)
    interpolate = create_interpolates(digit1, digit2, model)
    if index:
        interpolates = interpolate
        index = 0
    else:
        interpolates = torch.cat((interpolates, interpolate))
transform = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize(size=800),
    transforms.ToTensor()
])
interpolates = [transform(interpolates_) for interpolates_ in interpolates]
grid = utils.make_grid(interpolates, nrow = 9)
utils.save_image(grid, 'different_digit.png')
show(grid)

```